

Beyond Metadata for BBC iPlayer: an autoencoder-driven approach for embeddings generation in content similarity recommendation

Simone Spaccarotella

September 2024

Contents

1	Introduction and background	2
2	Outline of the issue or opportunity and the business problem to be solved	3
3	Scope of the project and Key Performance Indicators	5
4	Data selection, collection & pre-processing	5
5	Survey of potential alternatives	5
6	Implementation - performance metrics	5
7	Results	5
8	Discussion & conclusions/recommendations	5
9	Summary of findings	5
10	Implications	5
11	Caveats & limitations	5
12	Appendices	5

1 Introduction and background

I am a Software Engineer at the BBC, Team Lead for the Sounds web team, and I have been training as a Data Scientist, working in attachment with the iPlayer Recommendation team.

I built a machine learning model pipeline that produces content-to-content (C2C) similarity recommendations of video on-demand (VOD), for the "More Like This" section on BBC iPlayer [1]. This project is relevant to me because it is about recommendations, and I have been crossing paths with this world multiple times during my career at the BBC. I had a tangent encounter back in 2015, while working for a team that was building an initial recommender for BBC News and an API to provide recommendations using 3rd party engines. During a Hack Day some time later, I produced and presented a talk called "Recommendation Assumptions" [4], which was about recommendations and external factors affecting them, contextual to the consumption of the content itself.

The BBC is a well-known British broadcaster, and it is always evolving to remain relevant to its audience. Its mission is to inform, educate and entertain, and it operates within the boundaries set by the Royal Charter [2]. The current media landscape requires the BBC to deliver digital-first content that is relevant to the audience, and this involves investments in data and personalised services, not to mention a certain revolution in machine learning that is keeping everyone busy.

2 Outline of the issue or opportunity and the business problem to be solved

The BBC produces and stores a vast amount of metadata for its content, and this metadata is surfaced by countless services and APIs. One of the priority for the BBC is to increase the adoption across the business of "Passport" [3], an internal service that generates, stores, manages and provides access to a richer dataset of metadata annotations for multi modal content (audio, video and text). The usage in production is very low if not existent, and its adoption would make the access to metadata consistent BBC-wide, removing duplications and reducing effort and costs.

Furthermore, the similarity score of the current C2C recommender is directly proportional to the number of values in common between any pairs of items on a per-feature basis. But the commonality is calculated with an exact string equality, hence it ignores any relationship between different categorical values expressing a similar concept (e.g. "comedy", "stand-up comedy"). The number and types of tags are not enough to sufficiently describe the content. The data distribution is severely skewed towards the most popular category and no pre-processing is applied. Lastly, each similarity score is multiplied by a hardcoded weight that modulates the importance of a feature, but it doesn't solve the polarising effect of a skewed distribution. Unfortunately, because they are hyperparameters and not learned weights, the model can't improve its performances by minimising them against a cost function.

To address these issues, the aim of this project was:

- **To improve the quality of the C2C similarity recommendations.** The hypothesis was that by using in input a richer set of metadata that better describes the content, and by reducing the high-dimensional data to a lower-dimensional latent manifold, the model is able to generate embeddings that can improve the performances of the model by mapping the item similarity problem to a geometric distance calculation between vectors in a multi-dimensional Euclidean space.
- **To reduce the costs to generate C2C similarity recommendations.** I sourced the input data from Passport, to build a BBC-wide general solution that could be used by any product and applied to any type of content, because they all share the same set of annotations.
- **To build a foundational item-embeddings generator.** Content-based recommenders use information description to work. This project

had a twofold outcome: it created an immediate solution for C2C similarity that solely relies on item embeddings; it also created a foundational approach to generate embeddings for personalised recommender that combine user interactions with content metadata.

- 3 Scope of the project and Key Performance Indicators
- 4 Data selection, collection & pre-processing
- 5 Survey of potential alternatives
- 6 Implementation - performance metrics
- 7 Results
- 8 Discussion & conclusions/recommendations
- 9 Summary of findings
- 10 Implications
- 11 Caveats & limitations
- 12 Appendices

References

- [1] BBC. Bluey - "more like this" tab. <https://www.bbc.co.uk/iplayer/episodes/m000vbrk/bluey?seriesId=more-like-this>.
- [2] BBC. The Royal Charter. <https://www.bbc.com/aboutthebbc/governance/charter>, 2017. Valid until 31 December 2027.
- [3] BBC. How metadata will drive content discovery for the BBC online. <https://www.bbc.co.uk/webarchive/https%3A%2F%2Fwww.bbc.co.uk%2Fblogs%2Finternet%2Fentries%2Feachbb071-d471-4d85-ba9d-938c0c800d0b>, 2020. This page was archived on 1st August 2023 and is no longer updated.
- [4] Simone Spaccarotella. Recommendation assumptions. <https://www.slideshare.net/slideshow/recommendations-assumptions/236291920>, 2015. Presented at Prototyping Day @ Mozilla London on

3 September 2015, at Engineering Summit @ BBC on 7 March 2018,
uploaded on SlideShare on 27 June 2020.