

Beyond Metadata for BBC iPlayer:
an autoencoder-driven approach for
embeddings generation in content similarity
recommendation

Simone Spaccarotella

September 2024

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction and background | 3 |
| 2 | Outline of the issue or opportunity and the business problem to be solved | 4 |
| 3 | Methods and justification | 6 |
| 3.1 | Data pre-processing | 6 |
| 3.2 | Modelling and regularisation | 6 |
| 3.3 | Inference | 7 |
| 3.4 | Tools and frameworks | 7 |
| 4 | Scope of the project and Key Performance Indicators | 9 |
| 5 | Data selection, collection and pre-processing | 11 |
| 6 | Survey of potential alternatives | 13 |
| 7 | Implementation and performance metrics | 15 |
| 8 | Discussion and conclusions | 18 |
| 8.1 | Results | 18 |
| 8.2 | Summary of findings and recommendations | 18 |
| 8.3 | Implications | 19 |
| 8.4 | Caveats and limitations | 19 |

1 Introduction and background

I am a Software Engineer at the BBC, Team Lead for the Sounds web team, and I have been training as a Data Scientist, working in attachment with the iPlayer Recommendation team.

I built a machine learning model pipeline that generates content-to-content (C2C) similarity recommendations of video on-demand (VOD), for the “More Like This” section on BBC iPlayer [3]. This project is relevant to me because I have been crossing paths with the world of recommendations multiple times during my career at the BBC, and it sparked an interest in it. I had a tangent encounter back in 2015 while working for a team that was building an initial recommender for BBC News, and an API to provide recommendations using 3rd party engines. I also produced and presented a talk for a Hack Day. The talk was called “Recommendation Assumptions” [17], and it was about types of recommendations, and contextual external factors affecting them. Until now, when I was able to finally put my knowledge into practice with an actual project on real data.

The BBC is a well-known British broadcaster, and it is always evolving to remain relevant to its audience. Its mission is to inform, educate and entertain, and it operates within the boundaries set by the Royal Charter [7]. The current media landscape requires the BBC to deliver digital-first content that is relevant to the audience, and this involves investments in data and personalised services, not to mention a certain revolution in generative machine learning modelling that is keeping everyone busy.

2 Outline of the issue or opportunity and the business problem to be solved

The BBC produces and stores a vast amount of data for its content, and this data is produced and surfaced by countless services and APIs. One of the priority for the BBC is to increase the usage across the business of **Passport** [8], an internal BBC system that provides a richer set of metadata annotations for multi modal content (audio, video and text). The usage in production is very low, and its BBC-wide adoption would make the access to metadata consistent, removing duplications and reducing effort and costs.

Furthermore, the similarity score of the current C2C recommender is directly proportional to the number of values in common between any pairs of items on a per-feature basis. But the commonality is calculated with an exact string equality, ignoring any relationship between different categorical values expressing a similar concept (e.g. “comedy” and “stand-up comedy”).

Moreover, the number and types of tags are not enough to sufficiently describe the content, while the data distribution is severely skewed towards the most popular annotations, and no pre-processing is applied.

Lastly, each similarity score is multiplied by a hardcoded weight that modulates the importance of a feature, but it doesn’t solve the polarising effect of a skewed distribution. Unfortunately, because these are hyperparameters and not learned weights, the model can’t improve its performances by minimising them against a cost function.

To address these issues, the aim of this project was:

- **To improve the quality of the C2C similarity recommendations.** The hypothesis was that by using in input a richer set of metadata that better describes the content, and by reducing the high-dimensional data to a lower-dimensional latent manifold, the model would be able to generate embeddings that could improve the quality of the recommendations, by mapping the item similarity problem to a geometric distance calculation between vectors in a multi-dimensional Euclidean space.
- **To build a general solution that can be applied to multimodal content, reducing the costs.** I built a C2C solution for iPlayer, but I used Passport tags to make it general, so that it could be applied to any BBC content, because they all share the same set of common tags.
- **To build a foundational item-embeddings generator.** Content-based recommenders use item metadata. This project provided an

immediate solution for non-personalised C2C recommendations that solely relies on them. It also provided a foundational basis for personalised recommenders that could benefit from using content embeddings combined with other data like user interactions.

3 Methods and justification

3.1 Data pre-processing

I used **one-hot encoding** to transform the categorical features (i.e. the metadata annotations) into a numerical vector. It’s a simple yet effective encoding method and it’s perfect for the transformation of nominal categoricals, because it doesn’t introduce any ranking and/or arithmetic relationship among the encoded values. The downside of this approach is that it generates high-dimensional sparse arrays, introducing the so-called *curse of dimensionality* problem. Nonetheless, this was an accepted drawback that was managed in the modelling phase.

3.2 Modelling and regularisation

I trained an **autoencoder** [1, 15] to learn latent features of the Passport tags, and to reduce the size of the encoded vectors. The autoencoder is a type of encoder-decoder neural network, a self-supervised model capable of capturing non-linearity from the data. I used the “undercomplete” variant, that constrains the number of nodes in the hidden layers, creating a “bottleneck” of information flow through the network. This bottleneck is a form of *regularisation* that forces the model to learn latent attributes from the input, while reconstructing it with minimal loss. Ultimately, it prevents the model from *overfitting* the training dataset, by indexing it like a caching layer.

The trained *encoder* segment of the network, was used to compress the one-hot encoded high-dimensional sparse array into a lower-dimensional and denser representation called *embedding* [10]. This technique solved the curse of dimensionality and the data sparsity problems, and improved the calculation complexity and the quality of the recommendations at inference time.

To improve and assess the ability of the model to *generalise* on unseen data, I randomly shuffled the dataset and split it into 3 chunks: training, validation and test. I run **hyperparameter tuning** to find the best set of model parameters that minimised the objective function and I used the validation set to regularise the model with **early stopping**. This technique monitored the reconstruction loss on an out-of-sample dataset, allowing the model to stop training within a set “patience” threshold, after reaching a local minimum on the validation error.

I used **dropout** to further regularise the model and make it robust to small changes in the input, and **data augmentation**, by including in the training data the episodes of a programme that share the same tags with their parent container.

Weight decay and **batch normalisation** were tested during hyperparameter tuning and discarded for poor performances.

3.3 Inference

Item similarity was calculated with the **cosine** of the angle θ between each pair of embeddings [11]. This metric is insensitive to the magnitude of the vectors, and because high-frequency values tend to have a larger magnitude, it mitigates the impact of popularity in the similarity calculation. One-hot encoded vectors lack of meaningful relations between them. They represent unit vectors bound in the “positive quadrant” of a Cartesian coordinate system for a multi-dimensional Euclidean space. Because each pair can only have a finite number of angles, the cosine similarity will also assume a finite number of discrete values between 0 and 1, causing information loss. Ideally, we would expect the similarity score to assume a continuous value bound between -1 and 1, and this is only possible if the angle θ of any vector pair can assume a continuous value between 0 and 360 (i.e. 0π and 2π), hence the use of embeddings.

3.4 Tools and frameworks

The entire project was written in **Python**. It is the *de facto* programming language for data science and machine learning tasks. Python has an established, diverse and well-documented ecosystem of external libraries and frameworks that facilitated the job, and it is also the language of choice at the BBC.

I used **Pandas** only for tabular data manipulation, to generate and store the one-hot encoded vectors. Unfortunately, it wasn’t possible to use it for exploratory data analysis (EDA), because the iPlayer catalogue used in development had roughly one-year worth of data and it didn’t fit in memory, causing Pandas to crash. For this reason, I used **Dask**. It is a library capable of running out-of-memory and parallel execution for faster processing on single-node machines and distributed computing on multi-nodes machines, while using the familiar Pandas API.

I used **TensorFlow** and **Keras** for modelling, to build and train the encoder-decoder neural network architecture, and **Keras Tuner** for hyperparameters tuning. I also used **Scikit-learn** but not for modelling. It provided utility functions for the dataset splitting and the cosine similarity calculation, and I was already familiar with its API.

For data visualisation I used a combination of **Matplotlib** and **Seaborn**, while I used **rdflib** to fetch and parse the RDF documents from the BBC

Ontology. These documents represented the metadata values and contained the actual entity label. This label was needed to hydrate the set of metadata for each item, to visualise the recommendations for testing purposes. Worth also mentioning the use of **pytest** for unit testing and **black** for PEP 8 code compliance and formatting. Finally, I used **Jupyter Lab** to edit the project, **git** for code versioning, **GitHub** as a remote code repository and for collaboration, and **AWS Sagemaker** to run the pipeline on more capable virtual machines, especially during hyperparameter tuning.

4 Scope of the project and Key Performance Indicators

The scope for this project was to build an end-to-end machine learning solution, able to produce non-personalised content-to-content similarity recommendations, using Passport metadata tags as input.

The minimum viable outcome was to produce recommendations comparable with the ones currently in production. A desired outcome was to increase user engagement. The integration with Passport would make the solution general, and applicable to multimodal BBC content. If adopted, this solution would generate considerable savings.

I didn't have access to the actual cost for each AWS account used for C2C across the BBC, so I couldn't present a number to demonstrate the reduction. But, if this solution is adopted by N products for example, the final cost would be divided by a factor of N , which is a considerable improvement.

Comparability was a qualitative and subjective key performance indicator (KPI) that served as a compass, indicating whether the project was progressing towards the right direction. It was evaluated by several technical and non-technical stakeholders, with a diverse domain knowledge and background. I built a rudimentary visualisation tool that rendered the title, image, description and metadata of the seed programme and the top-K C2C recommendations. The people involved, could give their subjective feedback on what was their perceived level of similarity of the output, testing edge cases and common use cases with an expected outcome, sensitive recommendations (e.g. content recommendations for children accounts), and discussing anomalies and/or surprising results.

To measure user engagement though, the solution needed to be A/B tested in production, with real data. Unfortunately, there were too many moving parts outside my control that needed to happen, for me to even try and build a production-worthy version for A/B testing. Adopting this as a KPI would have delayed the project, increasing the odds of failure. To mitigate this risk, I had to decouple it from the success of the project.

I defined a hypothesis testable offline, within my area of control. The hypothesis stated that long-term user engagement is not just about accuracy. It can be affected by increasing diversity in the recommended content. A diverse set of recommendations generates new and unexpected results, with an increase in surprise and serendipity, pushing the user away from boredom. This was an untested hypothesis, but grounded in active research on the topic, such as [13] and [9], that supported the initial statement, and made it less far-fetched. For this reason, I defined a proxy metric that could

measure diversity offline, pending a future A/B testing for the validation of the hypothesis, which was pushed out of scope.

5 Data selection, collection and pre-processing

News and Sports articles, iPlayer videos on-demand, Sounds podcasts etc. are annotated with the so-called Passport tags. These tags describe the content produced by the BBC and can be used for retrieval (search) and filtering (recommendations). They can be applied either manually by an editorial team with domain knowledge, or semi-automatically by machine learning algorithms with human supervision.

Passport tags are distributed across the BBC via the universal content exposure and delivery (UCED) system, a self-service metadata delivery platform that exposes data as a document stream for products to integrate with. This platform provides different types of *consumers* such as REST API, AWS S3 bucket, etc. Passport documents are JSON objects that contain a property called “**taggings**”, an array of objects representing the metadata annotations. These objects in turn contain two properties: “**predicate**” and “**value**”. They represent the name and the value of a tag, and are expressed as URL-formatted strings, except for dates. The predicate is a class of the BBC Ontology [2] while the value can be a date or an entity defined as an RDF [20, 18] document, accessible in Turtle format [19] via the BBC Things API [4, 5, 6]. These entities are linked to each other and/or to external resources, and are described by attributes and relationships, giving the data a graph structure.

For the development of the project, I decided not to integrate with UCED but to use batches of Passport files, manually collected and stored on a local folder. This was a trade-off that allowed me to develop the project with real data while keeping the costs down, given that the resources needed to be set on two AWS accounts. Moreover, I didn’t want to pass the burden of maintenance to the team that owned the accounts, without having tested the feasibility of the solution first.

Content metadata is not classified as personally identifiable information (PII), according to the UK GDPR [12]. Nonetheless, this data is regulated by internal BBC data governance policies and as such, it is encrypted at rest and in transit by default. For this reason, no further precautions were required during storage and processing.

I chose to use Passport because this data provides a set of tags shared across all content produced by the BBC, making this a general solution that reduces duplications and ultimately costs. Passport provides a flexible and rich set of tags to describe any content. Annotations can describe canonical information such as *genre* and *format*, but also things like the *contributor* featuring in the programme, the *narrative theme*, the *editorial tone*, what the content is *about* or what relevant entities are *mentioned*, etc.

During pre-processing, a list of JSON files was loaded into the pipeline and the tags extrapolated into a dictionary data structure. The key of the dictionary was the programme ID (called *PID*) and the value was another dictionary describing the annotations. A programme can be tagged with the same predicate multiple times if it has different values, while the same value (e.g. “Music”) can be used by multiple predicates (e.g. **about** or **genre**).

The dictionary was transformed in a Pandas *Dataframe*, where the rows represented the programmes and the columns the tags. I used a MultiIndex [16] for the columns because I needed to keep the duplicate values (2nd-level index) across the predicates (1st-level index). I then populated the cells of the *Dataframe* with the value “1” if the programme was annotated with the corresponding tag, or “0” otherwise. This process generated one-hot encoded arrays. I initially started with a vectorisation approach known to be performant, using the “**get_dummies**” Pandas function. Surprisingly, it was slower and less scalable of the solution that I adopted in the end.

A source of bias in the dataset was the usage of the “**mentions**” tag. This tag is automatically generated by an algorithm that extracts terms deemed important, appearing in the text of an article or the transcript of an audio/video content. If something is “mentioned”, doesn’t necessarily describe what the content is about, because of the intrinsic ambiguities of natural languages. Figure of speech devices such as metaphors, analogies, allegories, etc., alter the meaning of a sentence for stylistic effect and can affect the representation of what the content was about. For example, if the phrase “being over the moon” is said referring to someone feeling great about something, unrelated to the cosmic space topic, and the term “moon” is extracted as a descriptor, it certainly increases the chances of misrepresentation. To mitigate this source of bias I dropped the tag in favour of “**about**”, a tag that describes what the content is really about. This tag is manually annotated by editorial teams who are trained to only annotate with relevant and topical entities.

A source of error was the encoding of unseen tags. I decided to drop the new tags and encode the data with the existing ones only. Also, a subset of programmes didn’t have any annotations. Creating an entry for them would have generated a minority of one-hot encoded arrays with all zeros, representing a characteristic class of uninformative observations. I decided to drop these programmes and include the only the ones with at least one annotation.

6 Survey of potential alternatives

Given a numerical vector representing the Passport tags of a programme in input, the model needed to return a list of the top-K programmes, sorted in a descending similarity order. The geometric interpretation of the similarity between two items, is the distance between the two vectors representing them.

I considered using a Clustering technique to group similar items together, so that given an item in input, the model would return the ones belonging to the same cluster. I didn't know how many clusters there could be in the data, and I didn't need to. I could have used a density-based technique to autodiscover them, but even in that case, some of the clusters could have had less than K items in a top-K similarity scenario. I could have returned the items belonging to the nearest cluster if needed, but at that point, what's was the meaning of clusters in that context anyway? In content similarity, any item has a degree of similarity with all the others and clustering was just a coarse-grained discretisation of that concept. I needed a more granular approach where every item could be compared with anyone else. In geometric terms it meant that I had to calculate the pairwise distance between any two pairs of vectors, so I didn't need to use clustering in the first place.

Both clustering and pairwise distance use the concept of geometrical distance between any two pairs, and this calculation is computationally expensive, and doesn't scale on big datasets, especially if it involves high-dimensional sparse arrays. Not only that, the one-hot encoding processing doesn't use any spatial proximity information to transform the categorical features into ones and zeros. It's a plain transformation that pivots the unique values of a given variable as features, and sets 1 to the corresponding annotations. If we take this raw vector as is, and interpret it as a multi-dimensional coordinate system and then calculate similarity, we are out of luck.

From a computational point of view, I needed to handle the curse of dimensionality, by reducing the dimension of the numerical vectors. Not only that, the new vectors needed to be in a denser and lower dimensional manifold embedded into the original high-dimensional ambient space, for the "local proximity" to have a meaning in terms of similarity.

Dimensionality reduction techniques such as Principal Component Analysis (PCA), Independent Component Analysis (ICA) or Linear Discriminant Analysis (LDA), were next on the list, but there was a problem with them too. Their job is to find a linear projection of the data, and this is a strong assumption that misses important non-linear structures. So, I discarded the idea of using PCA, but not the idea of using dimensionality reduction. I just needed a non-linear approach, so I turned my attention to manifold learning.

I'd like to point out that I also considered changing the pre-processing step to generate a different set of vectors that didn't suffer from the curse of dimensionality.

Hashing is a well-known non-invertible technique that can be used in machine learning for feature reduction. The reason why I didn't adopt this technique was because hashing has more hyperparameters than one-hot encoding, which increase its complexity. Settings like the size of the hash could have impacted the descriptive power of the vector even before starting the embedding learning phase. But most importantly, hashing functions introduce the collision problem, where two distinct entities can be mapped to the same index in the target domain. Because this is an issue that can't be removed, only mitigated, I tested the latest and strongest algorithm to reduce the likelihood of collisions, but the trade-off was too computationally expensive during training.

I talked about why I adopted one-hot encoding, let's now talk about the chosen approach for manifold learning, pros and cons and its justification.

7 Implementation and performance metrics

I used an undercomplete autoencoder to compress the high-dimensional one-hot encoded vectors to a lower dimensional embedding representation. This technique captured non-linearity by learning the underlying latent structure of the data, which was key to represent the original Passport tags in a geometrical space, exploiting local proximity as a measure of similarity.

Because the autoencoder has a symmetrical architecture between the *encoder* and the *decoder*, the input and output layers had the same dimensions: 8982 nodes. This number corresponded to the size of the one-hot encoded array.

Some of the hyperparameters were decided based on the nature of the problem. The input data was a tensor of zeros and ones, and the main objective of the network was to reconstruct the output with minimal loss. To achieve this, I used the *Sigmoid* as activation function for the output layer, and *binary cross-entropy* as loss function, to allow the network to push the reconstructed output values as close as possible to 0 and 1. I didn't use the *SoftMax* activation function for the output layer, because each single node needed to be able to assume a value between 0 and 1. I didn't use any residual-based loss function such as MSE or MAE, because the sum of the square (or absolute value) of the residuals for 0/1 values doesn't have a very steep slope to allow the gradient to move fast enough. Which is why the negative log loss in cross entropy is the best choice, because it massively penalises huge differences between \hat{y} and y with a logarithmic progression.

Some other hyperparameters were set after few rounds of training. For example, I used the *Rectified Linear Unit (ReLU)* as activation function for the hidden layers. I also tested other variants like *Leaky ReLU (LReLU)* and *Parametric ReLU (PReLU)* with some isolated hyperparameters tuning tests, but without any relevant improvements in performance. These hyperparameters were:

- the optimizer: Adam
- number of epochs: 100
- batch size: 300
- dropout rate: 0.2
- early stopping patience: 10
- early stopping monitoring: binary cross entropy on validation set

- data split: 80% training, 10% validation, 10% test

The remaining hyperparameters were selected through a final round of optimisations. I decided to use a "Bandit-based" approach called *Hyperband* [14], which improves upon *Random Search* by running fewer epochs on the randomly sampled set of parameters, and move on to the next stage by only testing the best performing ones, returning a ranked list of the best hyperparameter sets. The hyperparameters considered by Hyperband were:

- the number of hidden layers
- the embedding size
- the learning rate
- whether to use dropout
- whether to use batch normalisation

The number of hidden layers was a positive integer $N > 0$ where $N - 1$ layers were used for the encoder and the same amount for the decoder - because the autoencoder has a symmetrical structure - and one layer was instead used for the bottleneck in the middle. If the hyperparameters was set to 1, the network would only have the bottleneck.

The number of nodes per layer was also a positive integer N . Unfortunately, the number of combinations was too big to be meaningfully optimised. To reduce the complexity, I decided to couple the number of nodes per layer as a progression of integer divisions by 2. Each layer in the encoder had half the number of nodes compared to the previous one and double the amount of the successive one (and *vice versa* in the decoder). For example, because I had 8982 nodes in input, the progression of layer dimensions was: 4491, 2245, 1122, 561, 280, 140 and so on. Therefore, the embedding size was also bound to assume one of these values, greatly reducing the number of choices, hence the runtime of the hyperparameter optimisation.

The strengths of this approach were the efficient use of space and the fast inference time. Training took longer during hyperparameter tuning. This drawback was compensated by the fact that I adopted an offline training approach, together with a "stale-while-revalidate" policy to return the recommendations at inference time, in addition to the similarity scores being cached. The space scaled linearly with respect to the size of the input to store the embeddings, and quadratically to store the similarity scores. The drawback was the model's interpretability, or in this case, the lack thereof. The autoencoder is a *black box* by definition and the use of embeddings to

calculate the similarity just made the problem worse. It is difficult if not impossible to interpret which of the tags influenced the ranking in the top-K list by untangling the weights and biases of the neural network. But it's also difficult to explain what is going on, even using model-agnostic techniques like SHAP. The autoencoder's output is the reconstruction of the input, but the actual output of the entire model are the cosines of each pairs of vector embeddings. So, no explainability either, or at least, localised to the neural network.

To mitigate this problem, I considered "the model" being the entirety of the pipeline. Given a programme in input which is annotated with a set of tags, I get in output a list of the top-K most similar programmes, which are also annotated with a set of tags. This meant that I could analyse the tags and use their frequency and composition as proxy metrics for feature importance. I could calculate the intersection between the input tags and the recommended ones, giving a measure of coverage that can approximate the strength of the annotations of the input programme, conditioning the recommendations. I could also calculate the difference between the recommended tags minus the one in input, to calculate the distribution of the new tags introduced by the recommended programmes, to approximate diversity.

For this reason, and because the Sigmoid function can lead to loss saturation (i.e. the function plateau at the \pm infinity) that could prevent gradient-based learning algorithms from making progress, I'm using Binary Cross Entropy as a loss function. Binary Cross Entropy not only helps pushing the values closer to 0 or to 1 depending on the input value (to minimise the reconstruction loss), but also counteracts the effects of the exponential in the Sigmoid with the use of the log.

8 Discussion and conclusions

8.1 Results

This is a general solution that works on any content that uses Passport tags. It could be implemented as a single implementation to serve multiple products. This will reduce effort, duplication of code and data, and as a consequence, costs. I presented this project explaining the main intuition and showing the results with the visualisation tool. Once to the data scientists and engineers of the team I worked with and on a second occasion, to the stakeholders of the team that - after an internal restructuring - will be in charge to provide non-personalised recommendations, including C2C similarity. The feedback was positive in both cases. When I presented the project to the last team, it was mentioned that we would need to implement this solution so that we can assess end-to-end feasibility and put it in a position to be A/B tested against the current solution in production.

8.2 Summary of findings and recommendations

The results are perfectly aligned with the initial objectives and measure of success set at the beginning of the project. My recommendation is to build an initial minimum viable product (MVP) pipeline on the AWS development account. This will allow us to break down the engineering effort, and to spot any blockers/challenges that need to be addressed as early as possible, so that we can correct them and/or reconsider some of the assumptions ahead of the production build. We could use the Sagemaker Pipeline to build the stages, and initially test the end-to-end solution with batch Passport data. We would need two pipelines, one for training and one for inference, to generate the embeddings and the similarity scores. The output could be cached, to improve performances, and if everything goes smoothly from an engineering point of view, we could integrate with UCED to fetch real-time data, and prep the solution for A/B testing.

If this solution is viable and passes the A/B test, it could also be used to generate embeddings for other personalised recommenders that use item metadata in conjunction with user interactions and/or contextual data such as day and time of interaction, location, device used, etc.

The project could be further expanded by exploiting the graph nature of the data using graph neural network (GNN), specifically a graph autoencoder (GAE) to learn meaningful representation of the graph data, capturing the topological structure and the node content. This could improve upon the current autoencoder, that relies on the same data being flattened. This

effort will require further research and prototyping.

8.3 Implications

8.4 Caveats and limitations

Data drift can cause loss of similarity information, requiring a re-training of the model. The limitation is when new programmes with new unseen tags are added to the catalogue. If we don't re-run the pipeline, the new tags won't be one-hot encoded, which means they will be dropped entirely from the embedding generation.

9 Appendices

References

- [1] Dor Bank, Noam Koenigstein, and Raja Giryes. Autoencoders. *CoRR*, abs/2003.05991, 2020.
- [2] BBC. BBC Ontologies. <https://www.bbc.co.uk/ontologies/>.
- [3] BBC. Bluey - "more like this" tab. <https://www.bbc.co.uk/iplayer/episodes/m000vbrk/bluey?seriesId=more-like-this>.
- [4] BBC. Things. <https://www.bbc.co.uk/things>.
- [5] BBC. Things - About. <https://www.bbc.co.uk/things/about>.
- [6] BBC. Things - API. <https://www.bbc.co.uk/things/api>.
- [7] BBC. The Royal Charter. <https://www.bbc.com/aboutthebbc/governance/charter>, 2017. Valid until 31 December 2027.
- [8] BBC. How metadata will drive content discovery for the bbc online. <https://www.bbc.co.uk/webarchive/https%3A%2F%2Fwww.bbc.co.uk%2Fblogs%2Finternet%2Fentries%2Feachbb071-d471-4d85-ba9d-938c0c800d0b>, 2020. This page was archived on 1st August 2023 and is no longer updated.
- [9] Tomislav Duricic, Dominik Kowald, Emanuel Lacic, and Elisabeth Lex. Beyond-accuracy: A review on diversity, serendipity and fairness in recommender systems based on graph neural networks, 2023.
- [10] Google. Embeddings. <https://developers.google.com/machine-learning/crash-course/embeddings>.
- [11] Google. Measuring similarity from embeddings. <https://developers.google.com/machine-learning/clustering/dnn-clustering/supervised-similarity>.
- [12] UK Government. UK GDPR. <https://www.legislation.gov.uk/eur/2016/679/contents>, 2016. Regulation (EU) 2016/679 of the European Parliament and of the Council.
- [13] Marius Kaminskis and Derek G. Bridge. Diversity, serendipity, novelty, and coverage. *ACM Transactions on Interactive Intelligent Systems (TiiS)*, 7:1 – 42, 2016.

- [14] Lisha Li, Kevin G. Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. Hyperband: A novel bandit-based approach to hyperparameter optimization. *CoRR*, abs/1603.06560, 2016.
- [15] Umberto Michelucci. An introduction to autoencoders. *CoRR*, abs/2201.03898, 2022.
- [16] pandas via NumFOCUS, Inc. MultiIndex / advanced indexing. https://pandas.pydata.org/docs/user_guide/advanced.html, 2024.
- [17] Simone Spaccarotella. Recommendation assumptions. <https://www.slideshare.net/slideshow/recommendations-assumptions/236291920>, 2015. Presented at Prototyping Day @ Mozilla London on 3 September 2015, at Engineering Summit @ BBC on 7 March 2018, uploaded on SlideShare on 27 June 2020.
- [18] W3C. RDF 1.1 Concepts and Abstract Syntax. <https://www.w3.org/TR/rdf11-concepts>, 2014. W3C Recommendation 25 February 2014.
- [19] W3C. RDF 1.1 Turtle - Terse RDF Triple Language. <https://www.w3.org/TR/turtle/>, 2014. W3C Recommendation 25 February 2014.
- [20] W3C. Resource Description Framework (RDF). <https://www.w3.org/RDF>, 2014. Publication date: 2014-02-25 (with a previous version published at: 2004-02-10).